

CMSC417 Spring 2016 Lecture #12 3/21/2016

Agenda

- ⇒ grades posted for p2 & midterm
 - midterms handed back Wed
- ⇒ p3 assigned by tomorrow
- ⇒ office hours moved a bit this week

- ⇒ BGP (cont'd)
 - quick review
 - valley-free routing
 - policy examples
 - joining BGP & IGP
- ⇒ Reliable Transmission
 - sequence #s
 - stop and wait
 - sliding windows

How BGP selects paths

each path has attributes

→ considered (loosely) in the following order

□ weight (local to router, does not propagate)

□ local preference / local pref (local to AS, allows you to prefer routes from certain ASes)

□ shortest AS path


□ route origin (EGP over IGP over incomplete)

□ lowest MED (multi-exit discriminator, lets you tell other ASes how you'd prefer they enter you)


□ prefer eBGP over iBGP

□ lowest IGP metric to BGP nexthop

□ address / ID-based tie-breakers



local pref



Multi-Exit Discriminator (MED)

⇒ let's you tell others to prefer a way to get to you



⇒ e.g., reach me via terrestrial routes instead of a satellite link

shortest AS path

⇒ ASes prepend their AS number to the path when they re-advertise

⇒ ASes can prepend themselves multiple times, why would they do this?

route origin

⇒ prefer IGP over EGP over INCOMPLETE

generated locally

learned from outside


sourced from some other source

next hop

⇒ needs to be explicitly stated b/c the next hop router may not be the BGP speaker

⇒ needs to be reachable using the IGP

Joining BGP with your IGP



⇒ assume A, D, E are BGP speakers

BGP table for AS

prefix	BGP NH
18.0/16	E
12.0/12	A
128.8/15	D
128.69/16	A

IGP table at A

dst	NH
A	A
C	C
D	C
E	C

⇒ join IGP & BGP to find paths

18.0/16 → C

12.0/12 → A

128.8/15 → C

128.69/16 → A

Reliable Transmission


Two approaches

- ⇒ ARQ (Automatic Repeat reQuest)
- ⇒ FEC (Forward Error Correction)

We'll focus on ARQ

Stop & Wait

- ⇒ send a message
- ⇒ wait for an acknowledgement (ack)
- ⇒ resend after a timeout



- ⇒ problem: you can't tell what frame/message an ack was talking about
 - you also might accidentally resend on a lost ack

⇒ soln: sequence #s

- label frames/messages w/ sequence #s

- increment by 1 on each

- acks say what seq # they're from

- rolling over seq #s? how many do you need?

CMSC 417 Spring 2016 Lecture #12 3/21/2016

Stop & Wait

⇒ only one message / frame in flight at a time

⇒ pipe is not full

e.g., 1.5 Mbps link at 50 ms latency RTT

⇒ assuming 1500 byte frames

$$\frac{1500 \text{ bytes}}{50 \text{ ms}} \cdot \frac{1000 \text{ ms}}{s} \cdot \frac{\text{Kilobit}}{128 \text{ bytes}} = \sim 235 \text{ Kbps}$$

much less than 1.5 Mbps

bandwidth-delay product matters
latency × BW

sliding windows

⇒ have a large space of seq #s for frames

⇒ each frame gets the next seq #


⇒ 3 variables at sender

SWS = sliding window size in seq #s

LAR = last ack received (a seq #)

LFS = last frame sent (a seq #)

invariant $LFS - LAR \leq SWS$



Sliding Window Sizes

⇒ big

- allows for handling high latency and/or high bw situations

⇒ small

- gentler on the network
- gentler on receivers, esp. resource constrained devices
- if latency is low (or bw is low) it doesn't hurt you

⇒ no universal SWS

- to "fill the pipe" you need

$$\text{SWS} \cdot \text{message-len} \Rightarrow \underbrace{\text{bandwidth} \cdot \text{RTT}}$$

called bandwidth-delay product

TCP manages the SWS to meet these needs and a few other things.